

Introduction to Machine Learning - Exercise 1

Due Date: November 9th, 2021

Yosi Shrem, Yael Segal, Roni Chernyak and Yossi Keshet

October 26th, 2021

1 ERM

1.1

As mentioned in the class, a learning algorithm receives as input a training set S sampled from an unknown distribution \mathcal{D} and labeled by some target function f . Since the learner does not know what \mathcal{D} and f are, we use a training set of examples, which acts as a snapshot of the world that is available to the learner. In ERM we would like to find a solution that works well on that data.

An axis aligned classifier in the plane is a classifier that assigns the value 1 to a point if and only if it is inside a certain rectangle. Formally, given real numbers $a_1 \leq b_1, a_2 \leq b_2$, define the classifier $h_{(a_1, b_1, a_2, b_2)}$ by

$$h_{(a_1, b_1, a_2, b_2)}(x_1, x_2) = \begin{cases} 1 & \text{if } a_1 \leq x_1 \leq b_1 \text{ and } a_2 \leq x_2 \leq b_2 \\ 0 & \text{otherwise} \end{cases}$$



Let A be the algorithm that returns the smallest rectangle enclosing all positive examples in the training set. Explain whether A is an ERM or not.

Note: We rely on the realizability assumption. In another words, we assume that there is a rectangle that classifies correctly all the data points.

1.2

Let \mathcal{H} be the hypothesis space of binary classifiers over a domain \mathcal{X} . Let \mathcal{D} be an unknown distribution over \mathcal{X} , and let f be the target hypothesis in \mathcal{H} . Denote $h \in \mathcal{H}$.

Let us define the *true error* of h as,

$$L_{\mathcal{D}}(h) = \mathbb{P}_{x \sim \mathcal{D}}[h(x) \neq f(x)]$$

Let us define the *empirical error* of h over the training set S as,

$$L_S(h) = \frac{1}{m} \sum_{i=1}^m \mathbb{1}_{[h(x_i) \neq f(x_i)]}$$

where m is the number of training examples.

Show that the expected value of $L_S(h)$ over the choice of S equals $L_{\mathcal{D}}(h)$, namely,

$$\mathbb{E}_{S \sim \mathcal{D}}[L_S(h)] = L_{\mathcal{D}}(h)$$

2 Image Compression

Guidelines

1. You are not allowed to use external packages other than `os`, `sys`, `numpy` and `matplotlib`.
2. In order to submit your solution please upload your files to Submit and check your inbox for the feedback mail.
3. Technical questions about this exercise should be asked at the course' piazza.
4. Private/Personal issues regarding the deadline should be directed to **Yosi shrem**.

In this part of the exercise we will use the k-means algorithm for image compression, i.e. you should implement the k-means algorithm on the **image pixels** and then replace each pixel by its centroid.

You should implement the k-means algorithm as described in class (in recitation 2 presentation). You can use the following python snippet for reading, normalizing and reshaping the image so it will be ready for training.

```
import matplotlib.pyplot as plt
import numpy as np
import sys

image_fname, centroids_fname, out_fname = sys.argv[1], sys.argv[2], sys.argv[3]
z = np.loadtxt(centroids_fname) #load centroids

orig_pixels = plt.imread(image_fname)
pixels = orig_pixels.astype(float)/255.
# Reshape the image(128x128x3) into an Nx3 matrix where N = number of pixels.
pixels = pixels.reshape(-1, 3)
```

The image and the centroids initialization will be provided to you as an argument to your program. The run command to your program should be:

```
$ python ex1.py <image_path> <centroids_init_path> <output_log_fname>
```

For example:

```
$ python ex1.py dog.jpeg cents1.txt out.txt
```

When displaying your compressed image you should get similar results to the following:



Figure 1: dog.jpeg




Figure 2: Results

Reproducibility. Originally, the initial centroids in k-means are randomly generated. For reproducible purposes we provided you with the centroids initialization. Moreover, In case when 2 centroids are evenly close to a certain point, the one with the lower index "wins". To evaluate your program outputs, we provided you with 2 examples - `cents1.txt` and `cents3.txt` for centroid initialization, and with `out1.txt` and `out3.txt` the requested outputs. Please note that given these pre-defined values, your sequence of centroid updates should be deterministic and not random in any way.

Your code should run for **20 iterations or until convergence**. We define convergence where all the centroids don't change. Your program should **create** a file named `<output_log_fname>(3rd arg to your program)`, consisting of your centroids after each centroid update.

For example, when using `cents1.txt`, the requested output should be:

```
[iter 0]:[0.1327 0.1135 0.1088],[0.6819 0.6071 0.5152]
[iter 1]:[0.1022 0.0879 0.0899],[0.6549 0.5801 0.4896]
...
[iter 7]:[0.0918 0.0793 0.0837],[0.6435 0.569 0.4796]
[iter 8]:[0.0918 0.0793 0.0837],[0.6435 0.569 0.4796]
```

use the following line to match your output to the requested format:
`f"[iter {iter}]:{'',''.join([str(i) for i in new_z])}"`

As you can see the algorithm converged and stopped(same centroids). For consistency and speedup purposes, **after each centroids update, use the built-in `round(4)` function on each dimension** - precision of 4 digits.

```
>>> i=np.array([0.123456,0.987654])
>>> i.round(4)
array([0.1235, 0.9877])
```

Figure 3: `round(4)`

Once submitted you will get a feedback email describing whether the expected output matches the requested format- check it and correct if needed. Part of your grade will consist of automatic checks - follow the format guidelines.

3 What to submit?

You should submit the following files:

- A `txt` file, named `details.txt` with your name and ID.
- A PDF file named `ex1.pdf` with your answers to 1.1 and 1.2.
- Python 3.6+ file named `ex1.py`. The main function writes the centroids updates to `<output_log_fname>` as explained above.
- A PDF report named `report.pdf` including the following plots: The average loss/cost value as a function of the iterations for $k = 2, 4, 8, 16$ (4 plots in total, x-axis iterations, y-axis avg/total loss). Explain shortly about your centroids initialization process. You can initialize the centroids as you wish in this part.

Overall : `ex1.py`, `ex1.pdf`, `details.txt` and `report.pdf`

Good Luck!